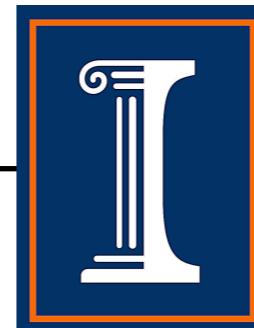


Uncertainty Quantification Group
at the University of Illinois at Urbana-Champaign

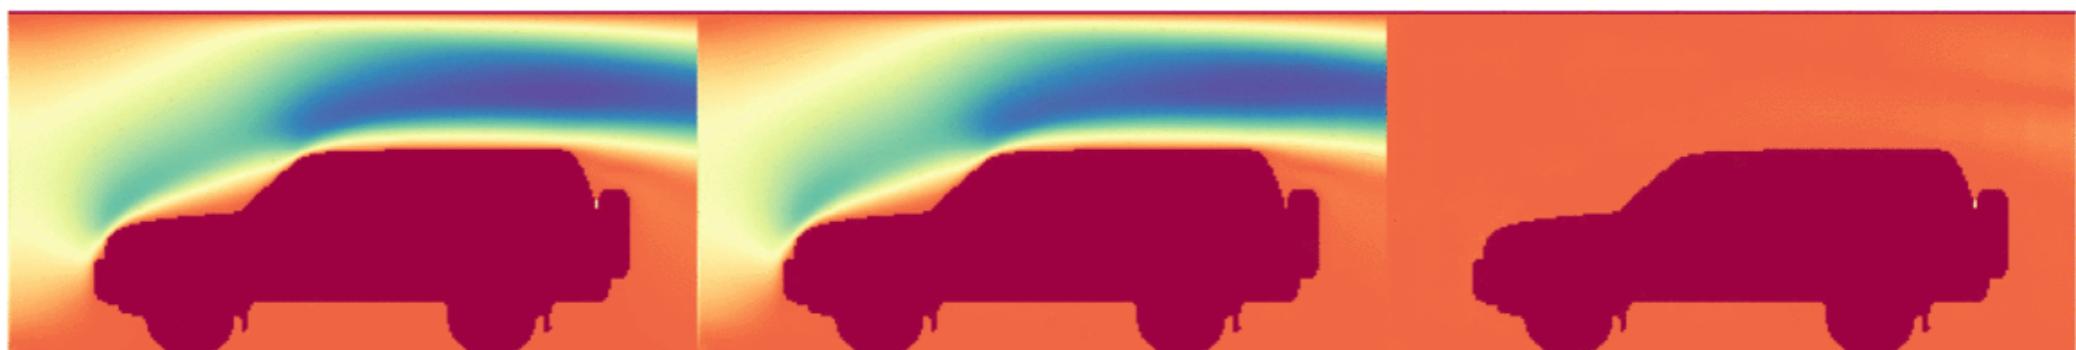


University of Illinois at Urbana-Champaign
Department of Civil and Environmental Engineering
UIUC UQ Group Presentation
Fall 2018

Physics-Informed Regularization of Deep Neural Networks

MOHAMMAD AMIN NABIAN

PhD Candidate
Sustainable and Resilient Infrastructure Systems Program



Contents

- Regression using deep neural networks
- Regularization of deep neural networks
- Physics-Informed (PI) regularization
- Numerical examples



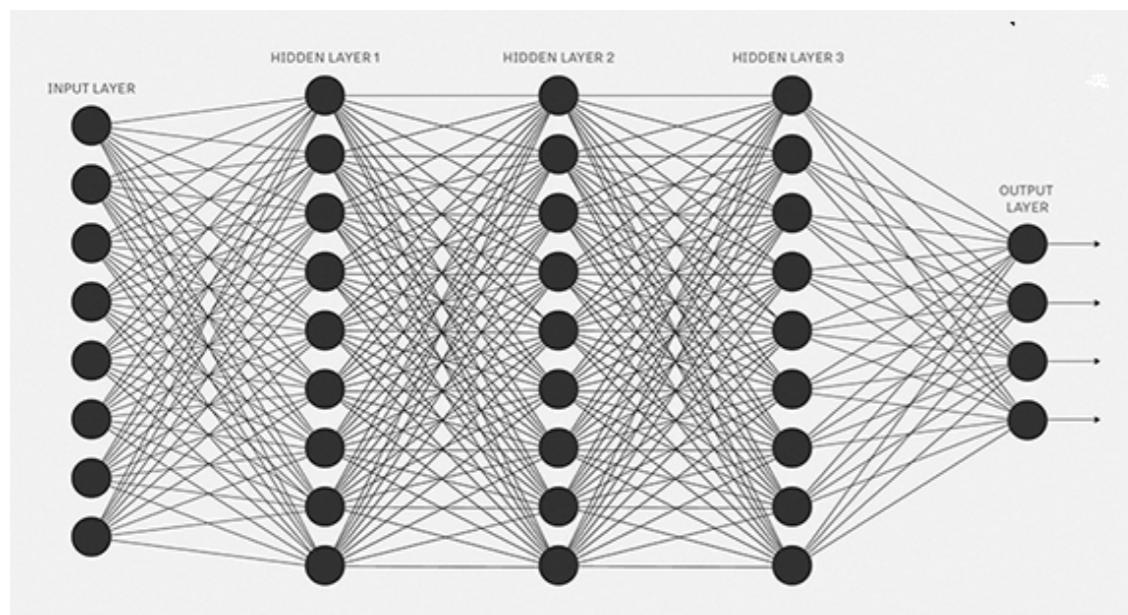
Regression Using Deep Neural Networks

Let us consider the following relationship holds:

$$\mathbf{y}_i = u(\mathbf{x}_i) + \epsilon_i$$

Our objective is to approximate this function by a deep neural network:

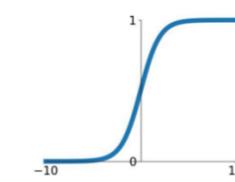
$$\mathbf{y}_i \approx \tilde{u}(\mathbf{x}_i; \Theta) = \sigma(\mathbf{x}_i \mathbf{W}_1 + \mathbf{b}_1) \mathbf{W}_2 + \mathbf{b}_2$$



Activation Functions

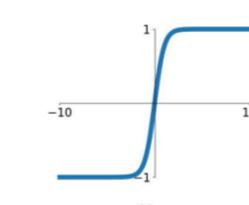
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



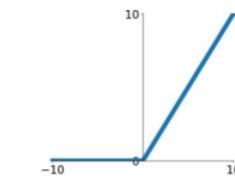
tanh

$$\tanh(x)$$



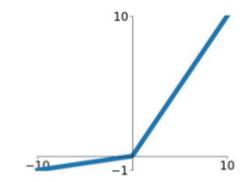
ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$

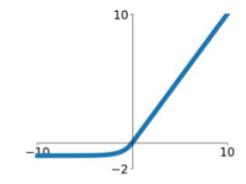


Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



In order to calibrate the weight matrices and biases, we may use Euclidean loss function as follows

$$J(\boldsymbol{\Theta}; \mathbf{X}, \mathbf{Y}) = \frac{1}{2n} \sum_{i=1}^n \|\mathbf{y}_i - \tilde{u}(\mathbf{x}_i; \boldsymbol{\Theta})\|^2$$

The model parameters can be calibrated according to the following optimization problem

$$\boldsymbol{\Theta}^* = \underset{\boldsymbol{\Theta}}{\operatorname{argmin}} J(\boldsymbol{\Theta}; \mathbf{X}, \mathbf{Y})$$

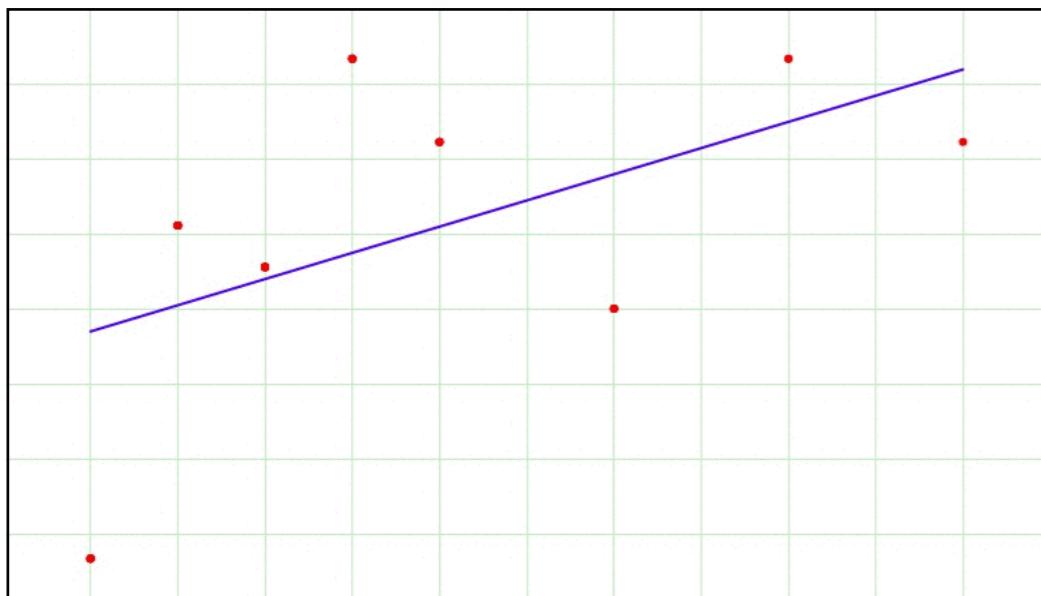
Using SGD, at the i^{th} iteration, the model parameters are updated according to

$$\boldsymbol{\Theta}^{(i+1)} = \boldsymbol{\Theta}^{(i)} - \eta^{(i)} \nabla_{\boldsymbol{\Theta}} J^{(i)}(\boldsymbol{\Theta}^{(i)}; \mathbf{X}, \mathbf{Y})$$



Regularization of Deep Neural Networks

Deep neural networks suffer from overfitting.



L1 Regularization

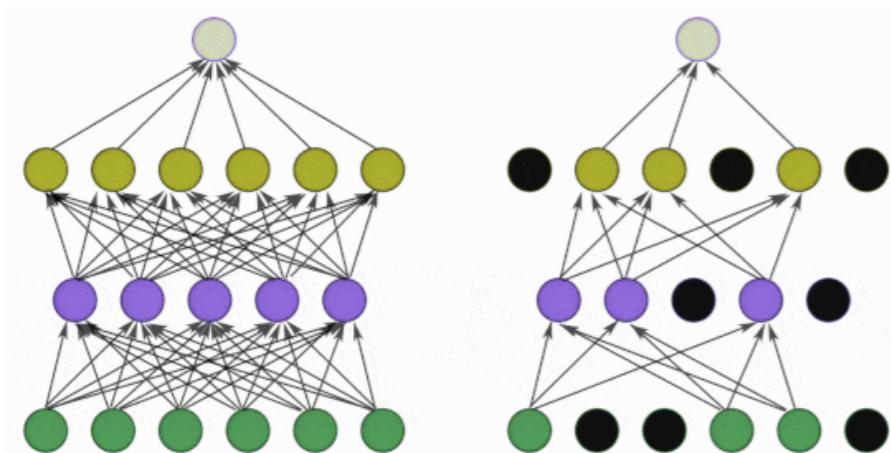
$$J_{L^1}(\Theta; \mathbf{X}, \mathbf{Y}) = J(\Theta; \mathbf{X}, \mathbf{Y}) + \lambda_1 \|\mathbf{W}\|_1$$

L2 Regularization

$$J_{L^2}(\Theta; \mathbf{X}, \mathbf{Y}) = J(\Theta; \mathbf{X}, \mathbf{Y}) + \frac{\lambda_2}{2} \mathbf{W}^T \mathbf{W}$$

Dropout

$$\mathbf{y} = \mathbf{r} \cdot \sigma(\mathbf{x}\mathbf{W}_1 + \mathbf{b}_1)\mathbf{W}_2 + \mathbf{b}_2$$



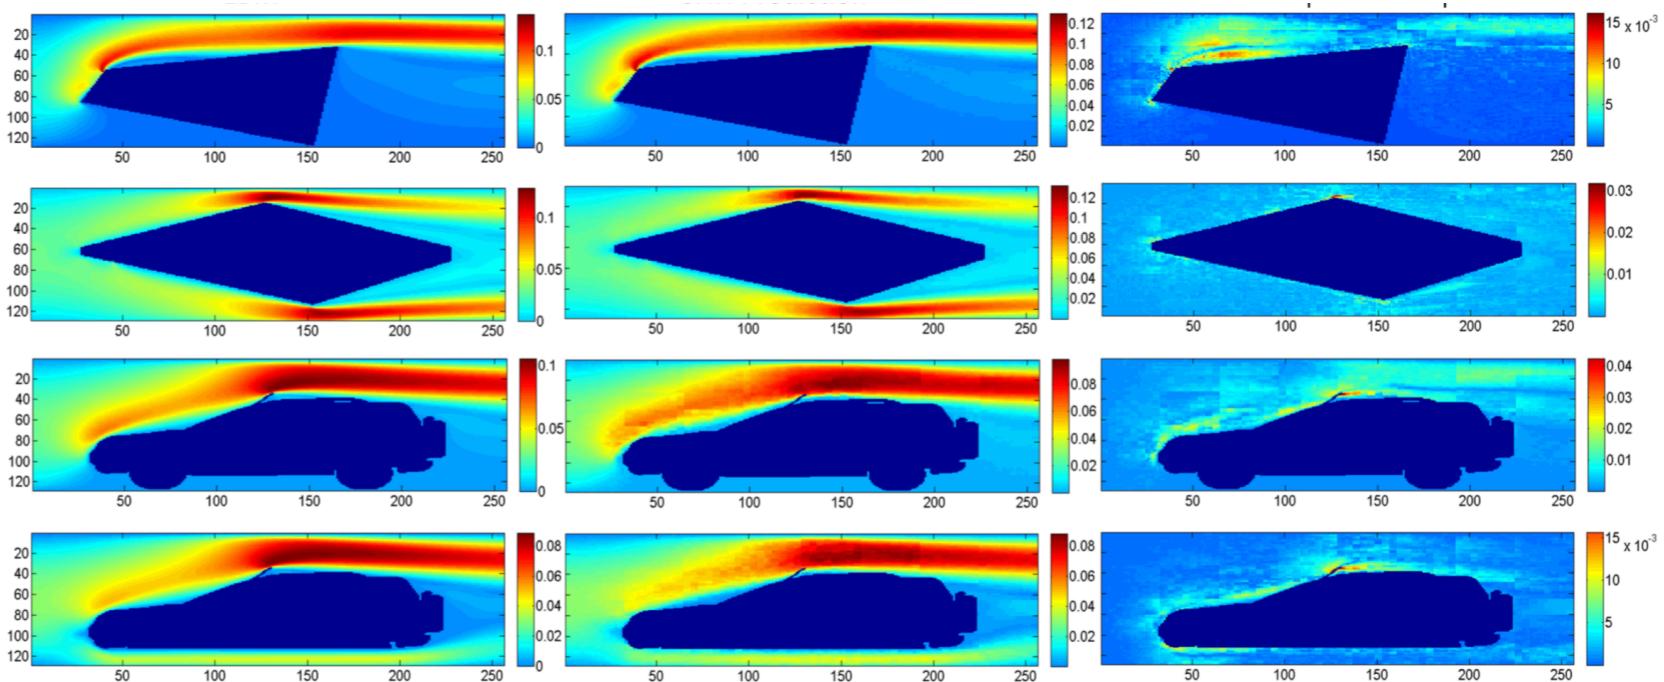
Physics-Informed (PI) Regularization

- Many science and engineering problems require repetitive simulation runs of a model with different input values.
- Examples are design optimization, model calibration, sensitivity analysis, what-if analysis, and design space exploration problems.
- However, in many real-world problems, obtaining a reliable outcome requires large number of these solves (typically for a partial differential equation), which can be prohibitive given the available resources.
- One way to alleviate this burden is to construct surrogate models that mimic the solution or response surface.

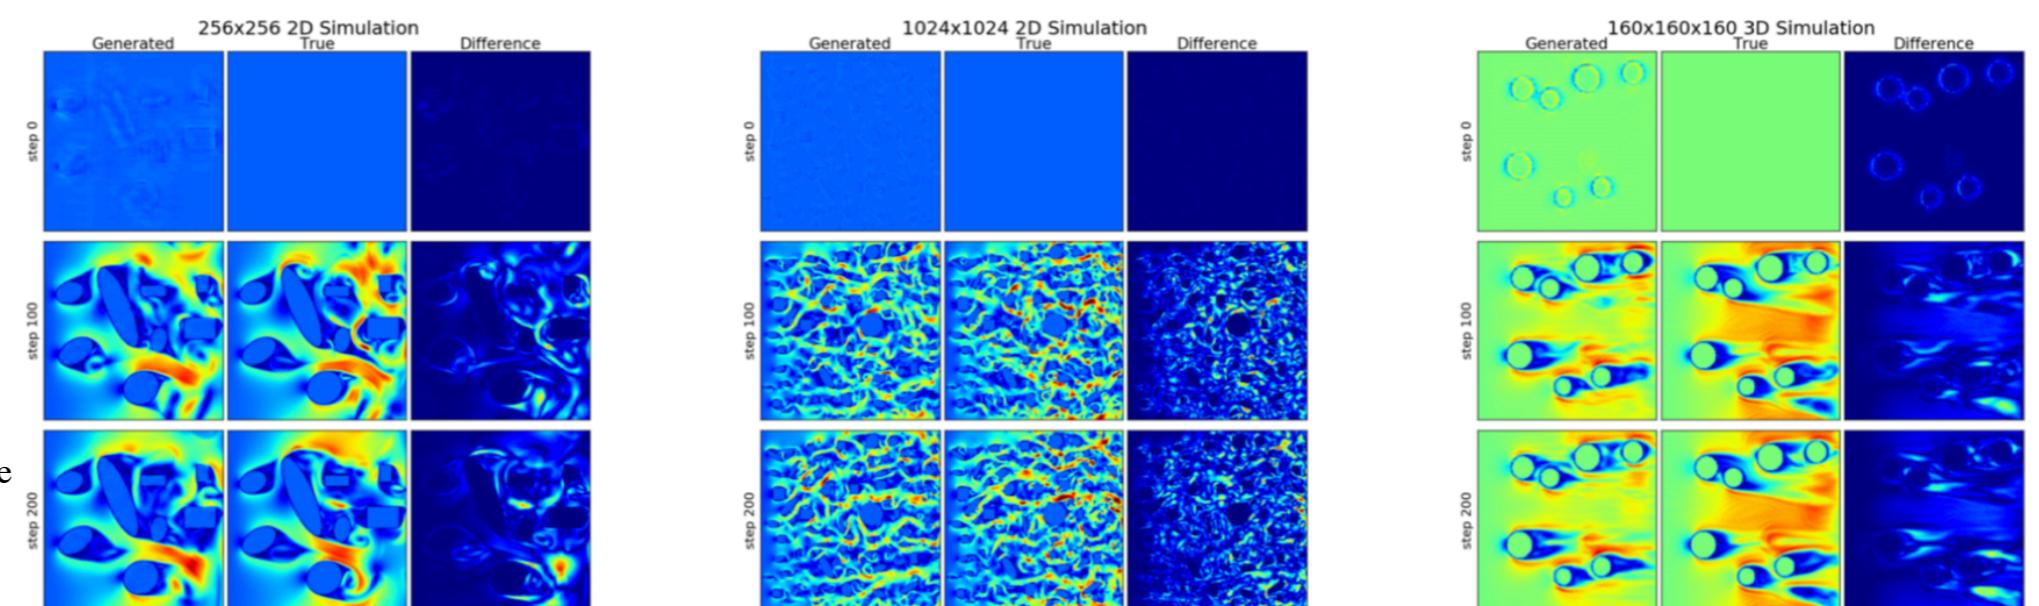


The state of the practice is to use pure data-driven regularization techniques, such as dropout, L2 and L1 regularizations.

Examples:



X. Guo, W. Li, F. Iorio, Convolutional neural networks for steady flow approximation, in: Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, ACM, 2016, pp. 481–490.



O. Hennigh, Lat-net: Compressing lattice boltzmann flow simulations using deep neural networks, arXiv preprint arXiv: 1705.09036.



- We presents a novel physics-informed approach for the regularization of deep neural network surrogates for systems that are subject to known governing laws which are in the form of a PDE.
- These governing equations are obtained using
 - First principles
 - Empirically-validated laws
 - Knowledge obtained by domain expertise
- In data-driven modeling of physical and biological systems, this prior knowledge is usually available, but not directly used in training of the models.



We consider cases where the response function a governing law, as follows

$$\mathcal{L}(\mathbf{x}_i, u(\mathbf{x}_i)) = 0, \quad \mathbf{x}_i \in \mathbb{R}^d.$$

The PI-regularized loss function is then defined as follows

$$J_{\text{PI}}(\boldsymbol{\Theta}; \mathbf{X}, \mathbf{Y}, \mathcal{L}) = J(\boldsymbol{\Theta}; \mathbf{X}, \mathbf{Y}) + \lambda_{\mathcal{L}} J_{\mathcal{L}}(\boldsymbol{\Theta}; \mathbf{X})$$

the physics-informed penalty term is defined as

$$J_{\mathcal{L}}(\boldsymbol{\Theta}; \mathbf{X}) = \frac{1}{2n} \sum_{i=1}^n [\mathcal{L}(\mathbf{x}_i, \tilde{u}(\mathbf{x}_i; \boldsymbol{\Theta}))]^2$$



By Adding the PI-regularization term to the standard loss function, the standard supervised learning task is converted to a **semi-supervised learning** task:

- The **supervised objective** minimizes the mean squared differences between model prediction and measurements.
- The **unsupervised objective** minimizes divergence from the governing laws.

The proposed PI regularization method effectively:

- Prevents deep neural networks from overfitting
- Results in surrogates that are better physically interpretable.

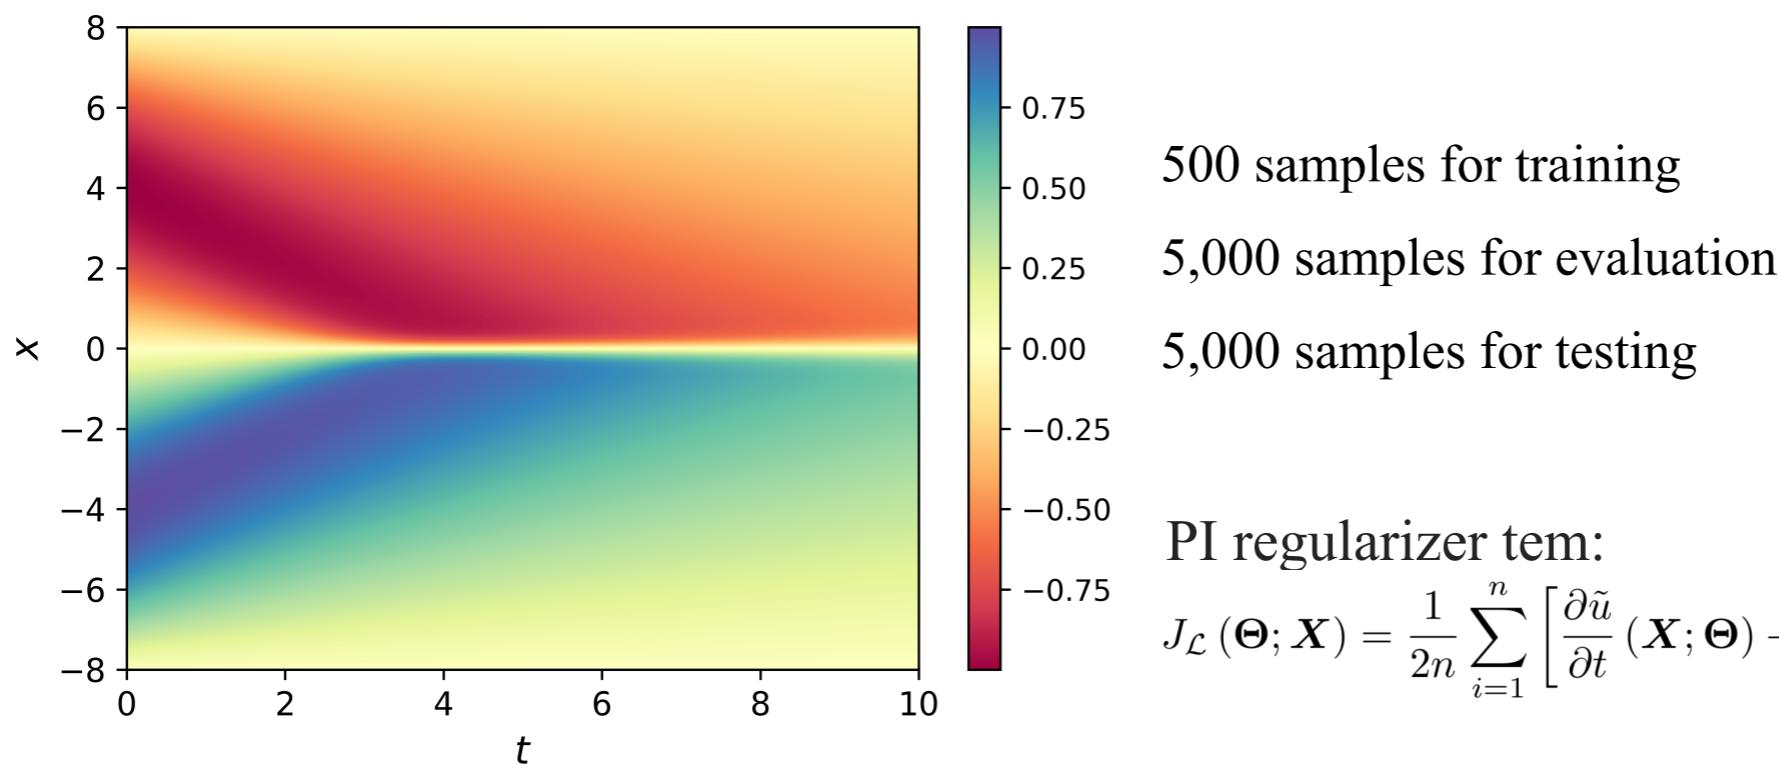


Example 1 : Burger's equation

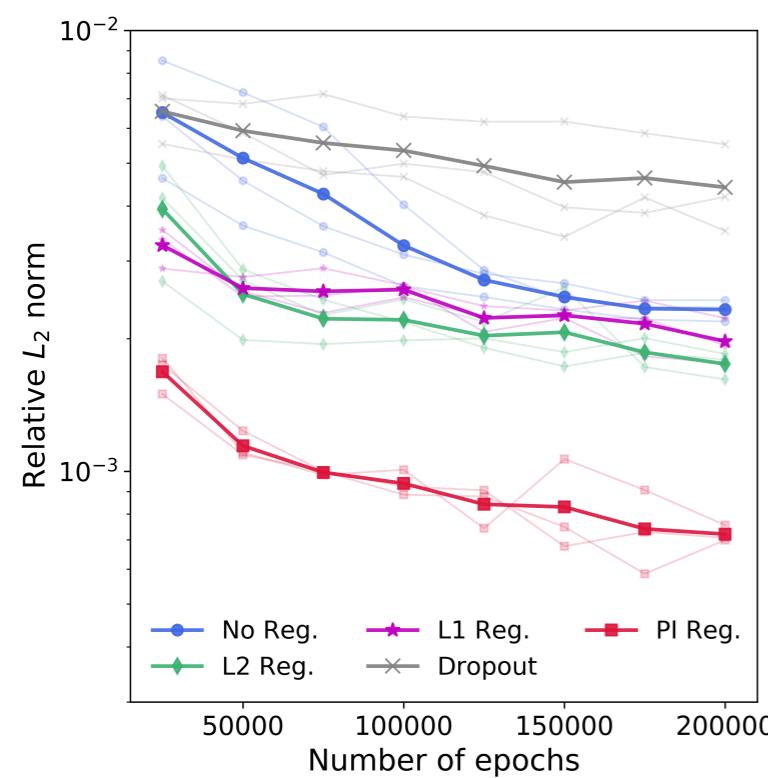
Let us start with the Burgers' equation, which arises in various areas of engineering, such as traffic flow, fluid mechanics :

$$\frac{\partial u}{\partial t} + u \frac{\partial u}{\partial x} - 0.1 \frac{\partial^2 u}{\partial x^2} = 0, \quad x \in [-8, 8], \quad t \in [0, 10],$$

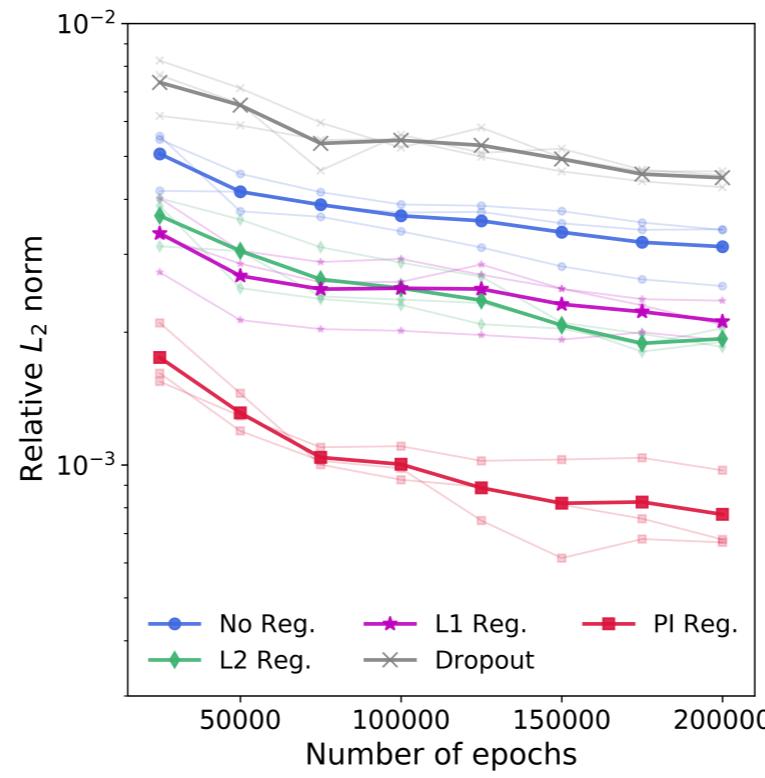
$$u(0, x) = -\sin(\pi x/8), \quad u(t, -8) = 0, \quad u(t, 8) = 0.$$



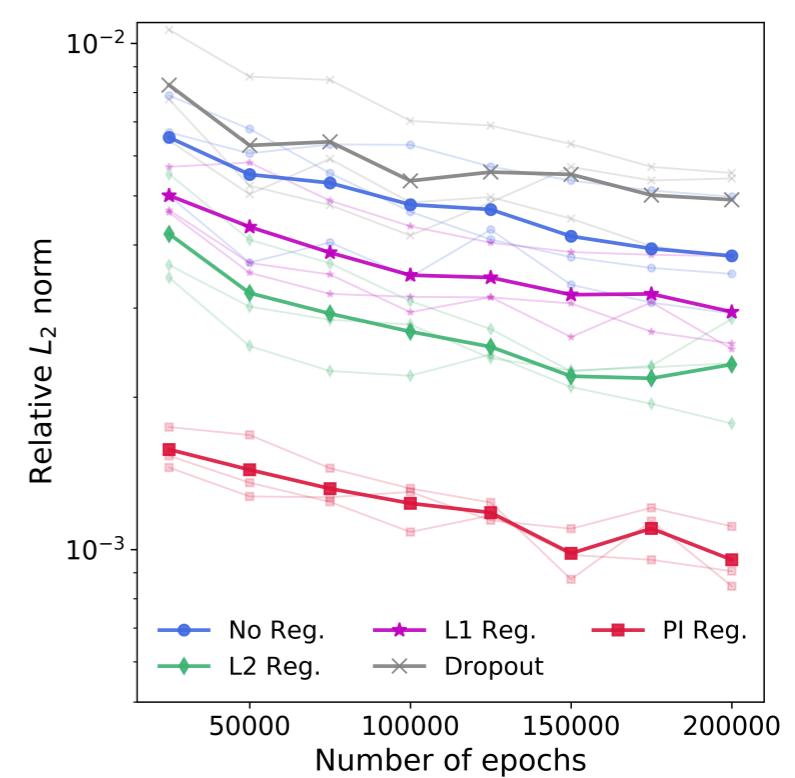
Comparison between the performance of different regularization methods



$$\gamma = 0$$

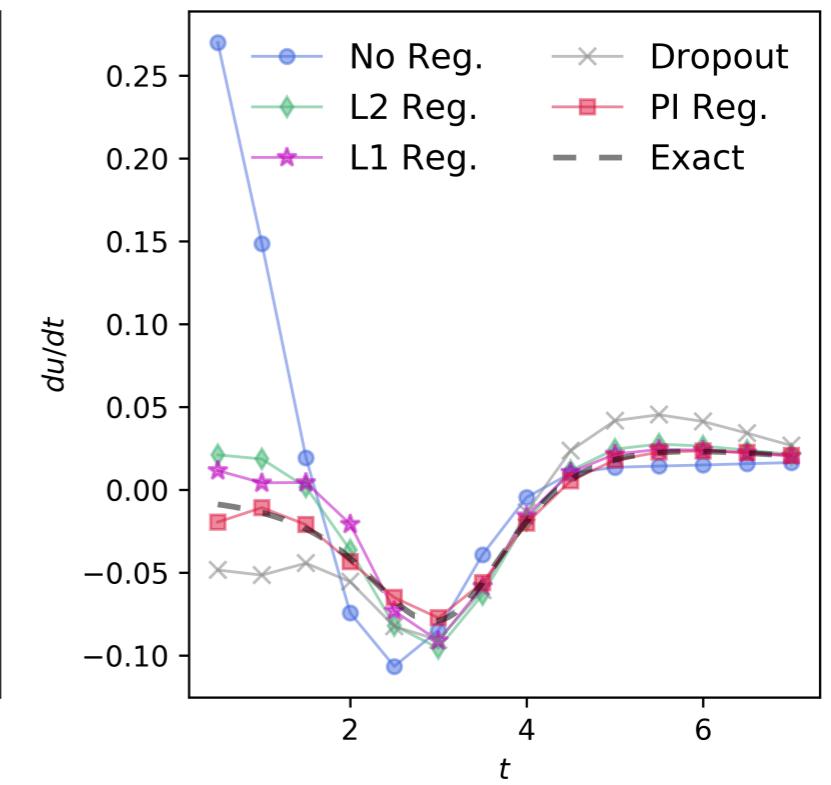
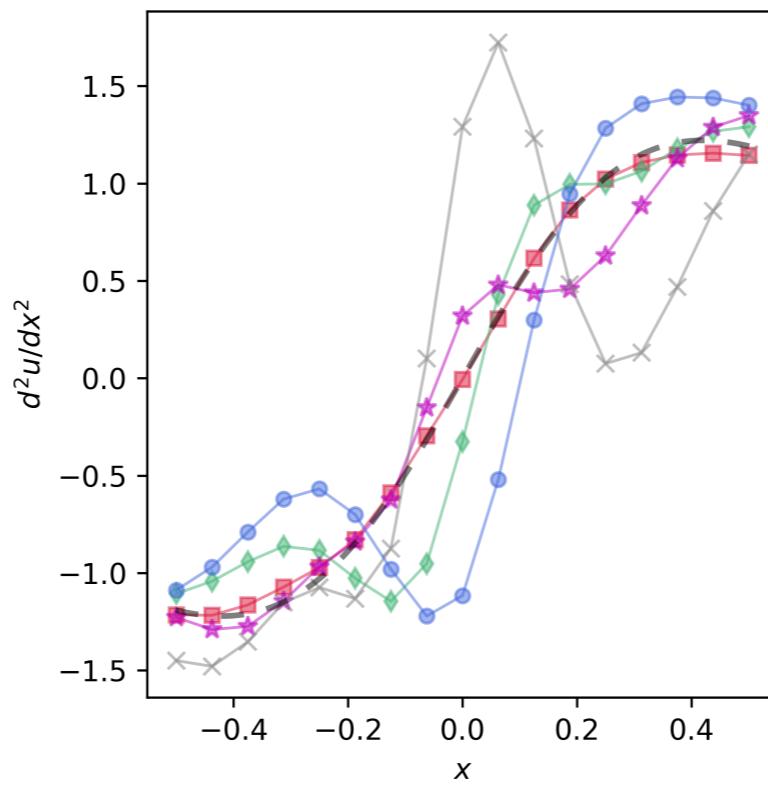
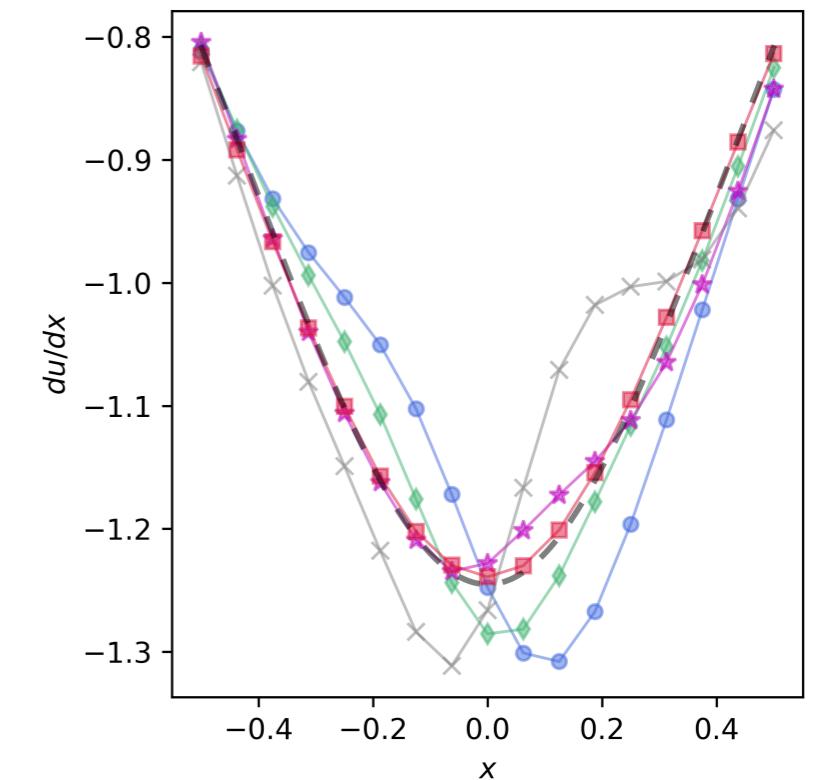
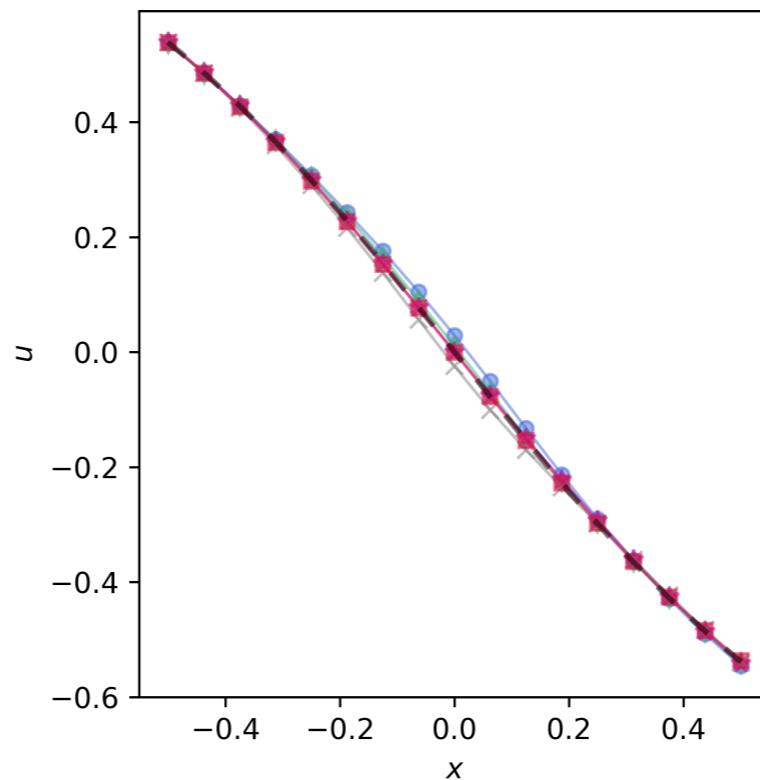


$$\gamma = 0.15$$



$$\gamma = 0.25$$





Comparison between the performance of different regularization methods in accurate prediction of first- and second-order derivatives



Example 2 : Navier-Stokes Equation

In this example we consider the vorticity equation given explicitly by

$$\frac{\partial \omega}{\partial t} = -u \frac{\partial \omega}{\partial x} - v \frac{\partial \omega}{\partial y} + 0.01 \left(\frac{\partial^2 \omega}{\partial x^2} + \frac{\partial^2 \omega}{\partial y^2} \right)$$

5,000 samples for training

15,000 samples for evaluation

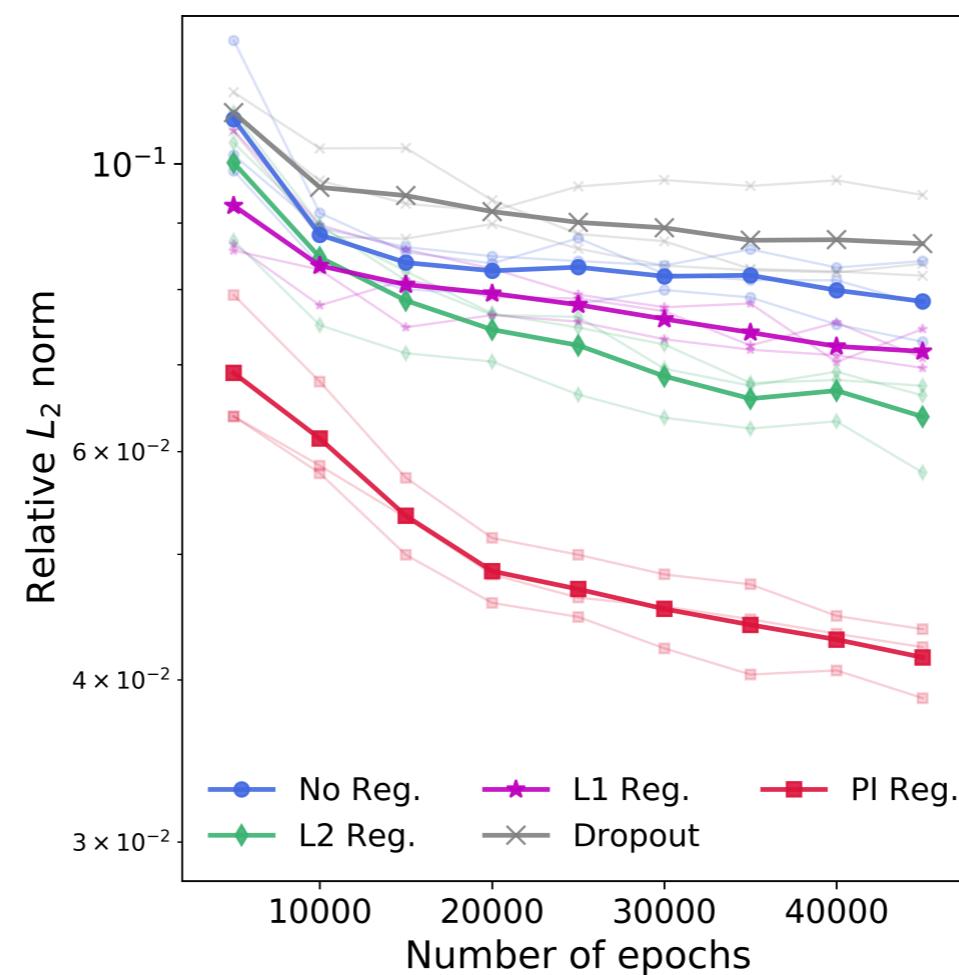
30,000 samples for testing

PI regularizer tem:

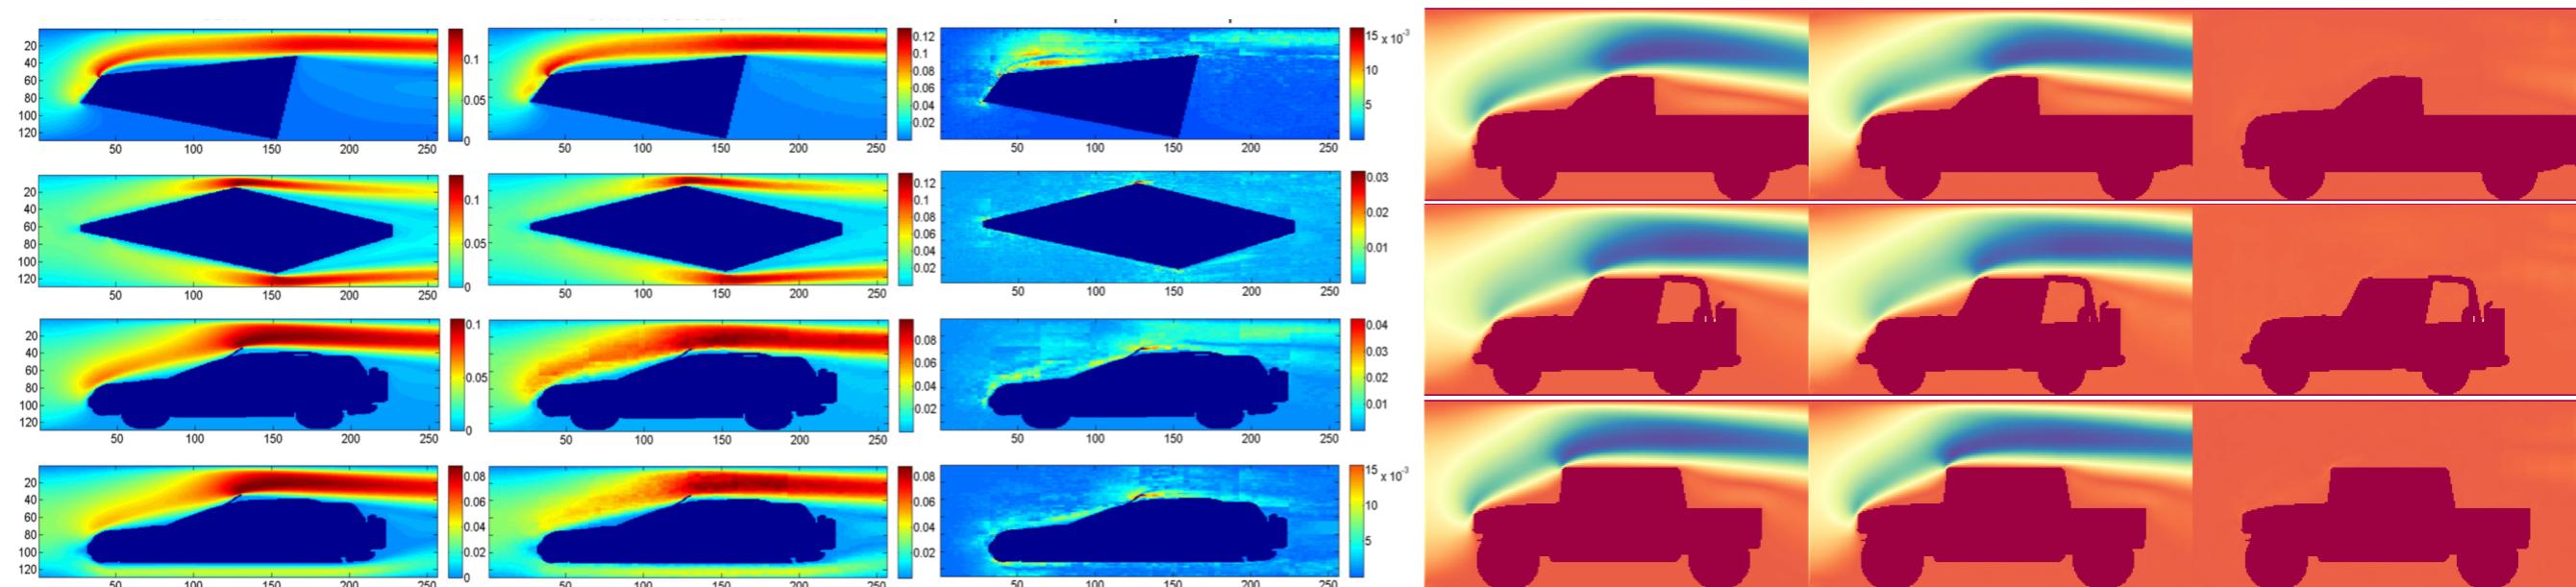
$$J_{\mathcal{L}} (\Theta; \mathbf{X}) = \frac{1}{2n} \sum_{i=1}^n \left[\frac{\partial \tilde{\omega}}{\partial t} (\mathbf{X}) + \tilde{u} (\mathbf{X}) \frac{\partial \tilde{\omega}}{\partial x} (\mathbf{X}) + \tilde{v} (\mathbf{X}) \frac{\partial \tilde{\omega}}{\partial y} (\mathbf{X}) - 0.01 \left(\frac{\partial^2 \tilde{\omega}}{\partial x^2} (\mathbf{X}) + \frac{\partial^2 \tilde{\omega}}{\partial y^2} (\mathbf{X}) \right) \right]^2$$



Comparison between the performance of different regularization methods



Example 3 : Surrogate modeling in CFD-based design optimization



Regularization method	No Reg.	Dropout ($P = 0.7$)	PI Reg. ($\lambda_{\mathcal{L}} = 0.1$)	Dropout ($P = 0.7$) & PI Reg. ($\lambda_{\mathcal{L}} = 0.1$)
Relative L_2 norm	1.41×10^{-5}	1.35×10^{-5}	1.36×10^{-5}	1.15×10^{-5}



Thank you!

